

INDEX

Practical No.	Title	Date	Page No.	Signature
01	REGRESSION MODEL Import a data from web storage. Name the dataset and now do Logistic Regression to find out relation between variables that are affecting the admission of a student in an institute based on his or her GRE score, GPA obtained and rank of the student. Also check the model is fit or not. require (foreign), require(MASS).		01	
02	MULTIPLE REGRESSION MODEL Apply multiple regressions, if data have a continuous independent variable. Apply on above dataset.		04	
03	CLASSIFICATION MODEL a. Install relevant package for classification. b. Choose classifier for classification problem. c. Evaluate the performance of classifier.		06	
04	CLUSTERING MODEL a. Clustering algorithms for unsupervised classification. b. Plot the cluster data using R visualizations.		09	
05	Install, configure and run Hadoop and HDFS ad explore		12	
06	Implement word count / frequency programs using MapReduce		17	
07	Implement Decision tree classification techniques		18	
08	Implement SVM classification techniques		21	

PRACTICAL NO: 01

AIM: REGRESSION MODEL Import a data from web storage. Name the dataset and now do Logistic Regression to find out relation between variables that are affecting the admission of a student in an institute based on his or her GRE score, GPA obtained and rank of the student. Also check the model is fit or not. require (foreign), require(MASS).

To perform Logistic Regression in R, we will use the UCLA Graduate School Admissions dataset, which is a standard dataset for this type of analysis.

Despite its name, **Logistic Regression** is used for **classification** (e.g., Admit vs. Not Admit). In R, we achieve this using the `glm()` function with a "binomial" family.

Step 1: Load Libraries and Import Data

We will use `foreign` to handle data structures and `MASS` for statistical functions.

How to execute: Open RStudio, go to `File > New File > R Script`, and paste the following:

```
R
# Load required libraries
require(foreign)
require(MASS)

# Import dataset from web storage
# This is a CSV file containing admit, gre, gpa, and rank
admission_data <- read.csv("https://stats.idre.ucla.edu/stat/data/binary.csv")

# View the first few rows
head(admission_data)
```

Step 2: Data Preparation

For a proper logistic regression, the `rank` of the school should be treated as a **factor** (categorical variable), not a continuous number.

```
R
# Convert rank to a factor
admission_data$rank <- factor(admission_data$rank)

# Summary of data to check for missing values
summary(admission_data)
```

Step 3: Run the Logistic Regression

We want to predict `admit` (the dependent variable) based on `gre`, `gpa`, and `rank` (the independent variables).

```
R
# Build the Logit model
# admit ~ . means predict admit using all other columns
logit_model <- glm(admit ~ gre + gpa + rank,
                     data = admission_data,
                     family = "binomial")

# View the results
summary(logit_model)
```

Note for the results:

- **P-values ($\Pr(>|z|)$):** If a p-value is less than **0.05**, that variable significantly affects admission.
- **Coefficients:** A positive estimate means an increase in that score increases the probability of admission.

Step 4: Evaluate Model Fit

In logistic regression, we don't use R^2 like in linear regression. Instead, we check the **Deviance** and use the **Likelihood Ratio Test** to see if the model is a good fit.

```
R
# 1. Check for Overall Model Significance
# Comparing our model to a 'null' model (a model with no predictors)
with(logit_model, pchisq(null.deviance - deviance, df.null - df.residual, lower.tail = FALSE))

# 2. Confusion Matrix to check accuracy
# We predict probabilities; if > 0.5, we classify as 'Admitted'
predictions <- predict(logit_model, type = "response")
predicted_class <- ifelse(predictions > 0.5, 1, 0)

# Create a table of Actual vs Predicted
table(Actual = admission_data$admit, Predicted = predicted_class)
```

Interpreting Fit:

- **Chi-Square Test:** If the result of the `pchisq` code above is very small (less than 0.05), your model fits significantly better than an empty model.
- **Deviance:** Lower Residual Deviance indicates a better fit.

OUT PUT:-

```
head(admission_data)
admit gre gpa rank
1     0 380 3.61    3
2     1 660 3.67    3
3     1 800 4.00    1
```

```
5      0 520 2.93      4
6      1 760 3.00      2
glm(formula = admit ~ gre + gpa + rank, family = "binomial",
     data = admission_data)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.989979	1.139951	-3.500	0.000465 ***
gre	0.002264	0.001094	2.070	0.038465 *
gpa	0.804038	0.331819	2.423	0.015388 *
rank2	-0.675443	0.316490	-2.134	0.032829 *
rank3	-1.340204	0.345306	-3.881	0.000104 ***
rank4	-1.551464	0.417832	-3.713	0.000205 ***

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 499.98 on 399 degrees of freedom
Residual deviance: 458.52 on 394 degrees of freedom
AIC: 470.52

Number of Fisher Scoring iterations: 4

```
[1] 7.578194e-08
>
> # 2. Confusion Matrix to check accuracy
  Predicted
Actual  0   1
  0 254 19
  1  97 30
```

PRACTICAL NO: 02

AIM : MULTIPLE REGRESSION MODEL Apply multiple regressions, if data have a continuous independent variable. Apply on above dataset.

In Multiple Linear Regression, we predict a **continuous** dependent variable. However, in the admissions dataset used previously, the variable `admit` is binary (0 or 1).

To apply **Multiple Regression** properly to this dataset, we will change our objective: We will predict the **GRE score** (a continuous variable) based on the student's **GPA** and the **Rank** of their school.

Step 1: Prepare the Environment

We will use the same dataset but treat `gre` as our target variable.

R

```
# Ensure libraries are loaded
library(MASS)

# Use the dataset from the previous step
# (Assuming admission_data is already loaded)
# If not: admission_data <- read.csv("https://stats.idre.ucla.edu/stat/data/binary.csv")

# Ensure rank is a factor
admission_data$rank <- factor(admission_data$rank)
```

Step 2: Build the Multiple Linear Regression Model

We use the `lm()` function (Linear Model) for continuous outcomes.

R

```
# Formula: gre ~ gpa + rank
# We are seeing how GPA and Rank together influence the GRE score
multiple_reg_model <- lm(gre ~ gpa + rank, data = admission_data)

# Display the statistical results
summary(multiple_reg_model)
```

Step 3: Interpret the Relation Between Variables

When you run the `summary()`, look for the following:

1. **Coefficients (Estimate):** This tells you the "slope." For example, if the GPA estimate is **50**, it means for every 1-point increase in GPA, the GRE score is expected to increase by 50 points, holding school rank constant.
2. **ZP-values (Pr(>|t|)):** If this is <0.05, the variable is a significant predictor of the GRE score.

3. **Adjusted R-squared:** This tells you what percentage of the variation in GRE scores is explained by your model.

Step 4: Evaluate Model Fit (Diagnostics)

To check if a linear model is "fit" for the data, we must check the residuals (the errors).

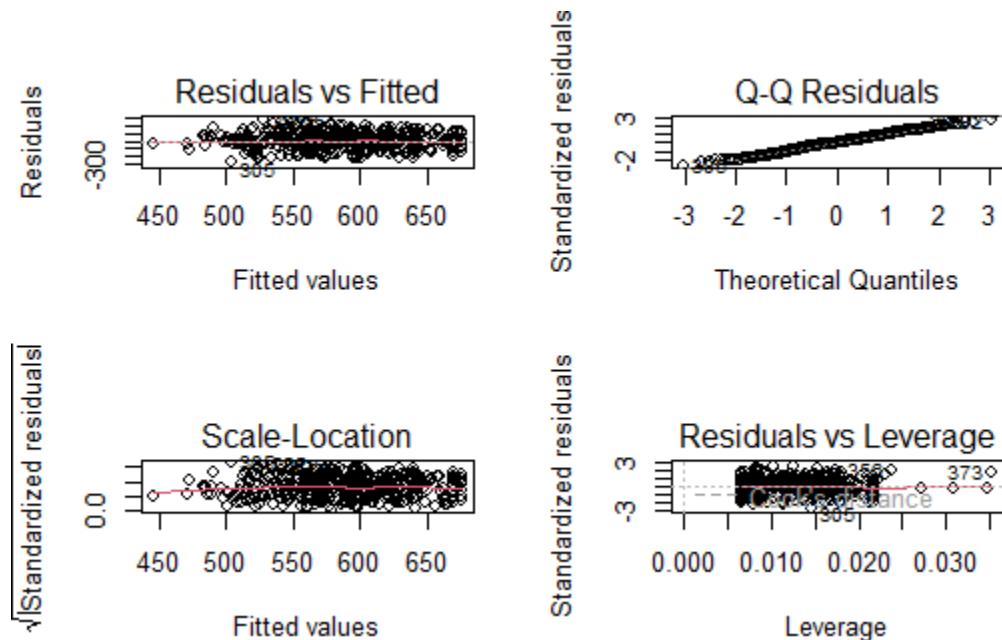
R

```
# Plot diagnostic graphs
# This will show 4 plots; press Enter in the console to see each one
par(mfrow = c(2, 2)) # Arrange plots in a 2x2 grid
plot(multiple_reg_model)
```

How to check if the model is fit:

- **Residuals vs Fitted:** You want to see a random scatter of points. If you see a curve, the relationship might not be linear.
- **Normal Q-Q:** The points should follow the diagonal dashed line. If they do, your errors are normally distributed (which is good).

OUTPUT:



PRACTICAL NO: 03

AIM: CLASSIFICATION MODEL a. Install relevant package for classification. b. Choose classifier for classification problem. c. Evaluate the performance of classifier.

Building a classification model in R is a streamlined process thanks to the `caret` (Classification And REgression Training) package, which provides a consistent interface for hundreds of different algorithms.

Here is a step-by-step guide to building, running, and evaluating a **Random Forest** classifier using the classic `iris` dataset.

Step 1: Install and Load Relevant Packages

First, you need to install the tools. The `caret` package is the industry standard for machine learning in R, and `randomForest` is the specific engine we will use.

How to execute: Type these lines into your R script or the Console and press **Enter**.

```
# Install packages (only need to do this once)
install.packages("caret")
install.packages("randomForest")
```

```
# Load the libraries
library(caret)
library(randomForest)
```

Step 2: Prepare the Data

Before choosing a classifier, we must split the data into a **Training set** (to teach the model) and a **Test set** (to see if it actually learned).

```
# Load built-in dataset
data(iris)

# Set a seed for reproducibility (so you get the same results every time)
set.seed(123)

# Create an index to split 80% for training, 20% for testing
trainIndex <- createDataPartition(iris$Species, p = 0.8, list = FALSE)
trainData <- iris[trainIndex, ]
testData <- iris[-trainIndex, ]
```

Step 3: Choose and Train the Classifier

We will use **Random Forest** because it is robust, handles non-linear data well, and is highly accurate for classification tasks.

R

```
# Train the model
# Species ~ . means "Predict Species using all other variables"
model <- train(Species ~.,
  data = trainData,
  method = "rf",
  trControl = trainControl(method = "cv", number = 5)) # 5-fold cross-validation

# View model details
print(model)
```

Step 4: Evaluate Performance

To evaluate the model, we use it to predict the species of the "unseen" test data and compare those predictions to the actual results using a **Confusion Matrix**.

R

```
# Make predictions on the test set
predictions <- predict(model, testData)

# Evaluate using a Confusion Matrix
evaluation <- confusionMatrix(predictions, testData$Species)

# Print the results
print(evaluation)
```

OUTPUT:

Confusion Matrix and Statistics

		Reference		
Prediction		setosa	versicolor	virginica
setosa	setosa	10	0	0
versicolor	versicolor	0	10	2
virginica	virginica	0	0	8

Overall statistics

```
Accuracy : 0.9333
95% CI : (0.7793, 0.9918)
No Information Rate : 0.3333
P-Value [Acc > NIR] : 8.747e-12
```

Kappa : 0.9

Mcnemar's Test P-Value : NA

Statistics by Class:

	Big Data Analytics		
	Class: setosa	Class: versicolor	Class: virginica
Sensitivity	1.0000	1.0000	0.8000
Specificity	1.0000	0.9000	1.0000
Pos Pred Value	1.0000	0.8333	1.0000
Neg Pred Value	1.0000	1.0000	0.9091
Prevalence	0.3333	0.3333	0.3333
Detection Rate	0.3333	0.3333	0.2667
Detection Prevalence	0.3333	0.4000	0.2667
Balanced Accuracy	1.0000	0.9500	0.9000

Understanding the Output

- **Accuracy:** The percentage of correct predictions. (e.g., 0.96 means 96% correct).
- **Kappa:** A measure of how well the classifier performed compared to how well it would have performed by simply guessing.
- **Confusion Matrix Table:** Shows exactly which species were misclassified (e.g., if a *Versicolor* was mistaken for a *Virginica*).

PRACTICAL NO: 04

AIM: CLUSTERING MODEL a. Clustering algorithms for unsupervised classification. b. Plot the cluster data using R visualizations.

Clustering is an unsupervised learning technique used to group data points based on similarities without using pre-defined labels. The most common algorithm for this is **K-Means Clustering**.

We will use the `iris` dataset again, but we will hide the "Species" column to see if the algorithm can discover the groupings on its own.

Step 1: Prepare the Data

Clustering requires numeric data. Since the `iris` dataset has 4 numeric columns and 1 categorical (Species), we will strip the species names so the model works "blind."

How to execute: Copy this into a new R Script and run it.

```
R
# Load the dataset
data(iris)

# Remove the Species column (column 5) for unsupervised learning
iris_cluster <- iris[, -5]

# Scale the data (Standardization)
# This ensures variables with larger numbers don't dominate the model
iris_scaled <- scale(iris_cluster)
```

Step 2: Choose the Number of Clusters (The "Elbow" Method)

Before running K-Means, we need to decide on **k** (the number of groups). We use the Elbow Method to find the point where adding more clusters doesn't significantly improve the fit.

```
R
# Determine number of clusters
set.seed(123)
wss <- sapply(1:10, function(k){kmeans(iris_scaled, k, nstart=20 )$tot.withinss})

# Plot the elbow curve
plot(1:10, wss, type="b", pch = 19, frame = FALSE,
      xlab="Number of clusters K",
      ylab="Total within-clusters sum of squares")
```

Look for the "bend" in the arm. For Iris, the bend is usually at $k = 3$.

Step 3: Apply the K-Means Algorithm

Now we apply the algorithm using 3 clusters.

R

```
# Apply K-Means
set.seed(123)
km_result <- kmeans(iris_scaled, centers = 3, nstart = 25)

# View the cluster assignments
print(km_result$cluster)
```

Step 4: Visualize the Clusters

Since we have 4 variables, we can't easily graph them in 2D. We use the `factoextra` package to perform Principal Component Analysis (PCA) automatically and plot the clusters clearly.

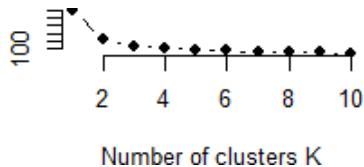
How to execute: You may need to install `factoextra` first.

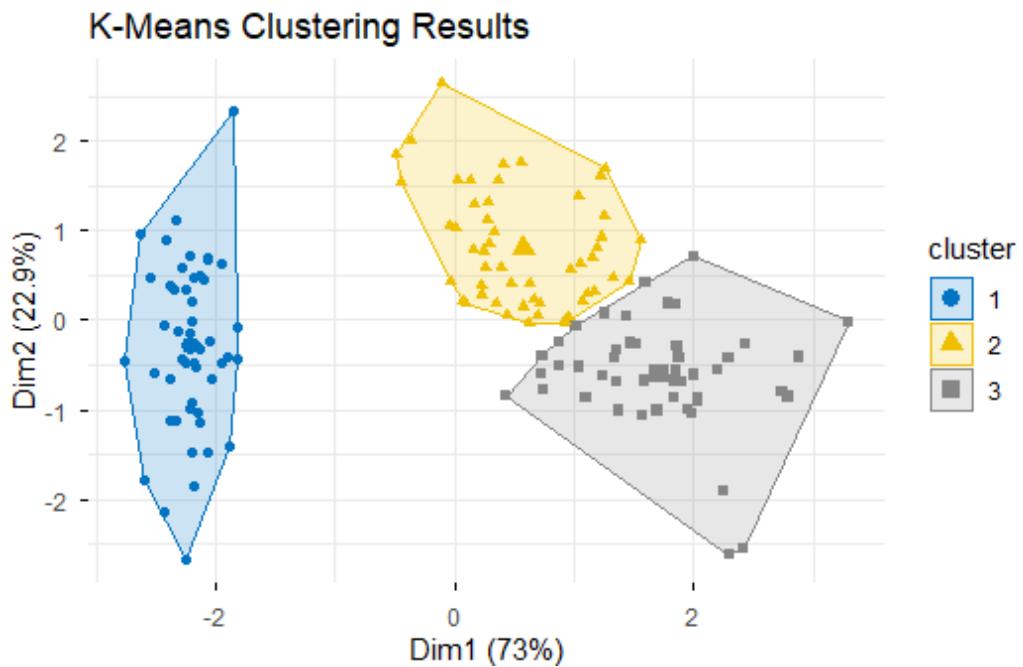
R

```
# Install and load visualization package
if(!require(factoextra)) install.packages("factoextra")
library(factoextra)

# Plot the clusters
fviz_cluster(km_result, data = iris_scaled,
             palette = "jco",
             geom = "point",
             ellipse.type = "convex",
             ggtheme = theme_minimal(),
             main = "K-Means Clustering Results")
```

OUTPUT:





PRACTICAL NO: 05**AIM: Install, configure and run Hadoop and HDFS ad explore.**

1. Install Java JDK 1.8
2. Download Hadoop and extract and place under C drive
3. Set Path in Environment Variables
4. Config files under Hadoop directory
5. Create folder datanode and namenode under data directory
6. Edit HDFS and YARN files
7. Set Java Home environment in Hadoop environment
8. Setup Complete. Test by executing start-all.cmd

There are two ways to install Hadoop, i.e.

9. Single node
10. Multi node

Here, we use multi node cluster.**1. Install Java**

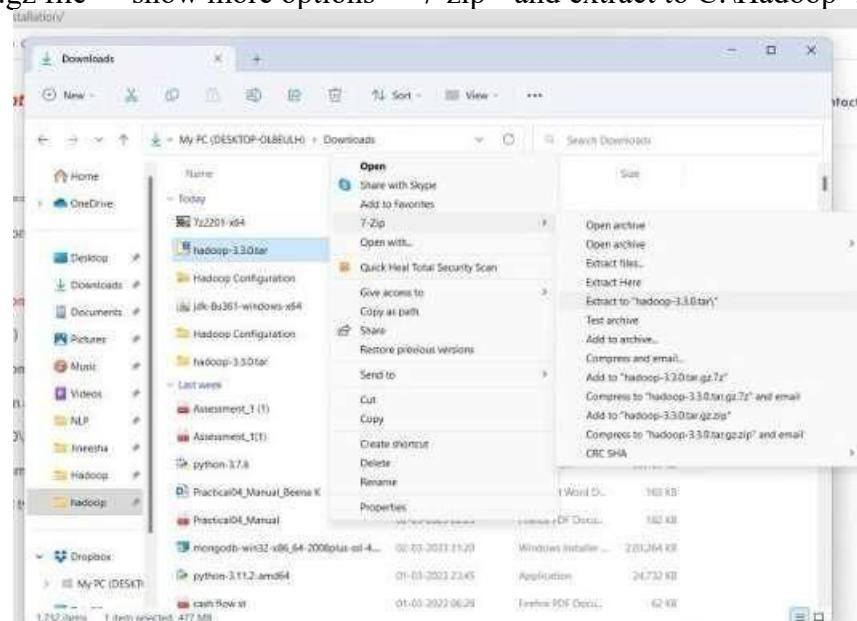
11. – Java JDK Link to download <https://www.oracle.com/java/technologies/javase-jdk8-downloads.html>
12. – extract and install Java in C:\Java
13. – open cmd and type -> javac -version

```
C:\Users>cd Beena

C:\Users\Beena>java -version
java version "1.8.0_361"
Java(TM) SE Runtime Environment (build 1.8.0_361-b09)
Java HotSpot(TM) 64-Bit Server VM (build 25.361-b09, mixed mode)
```

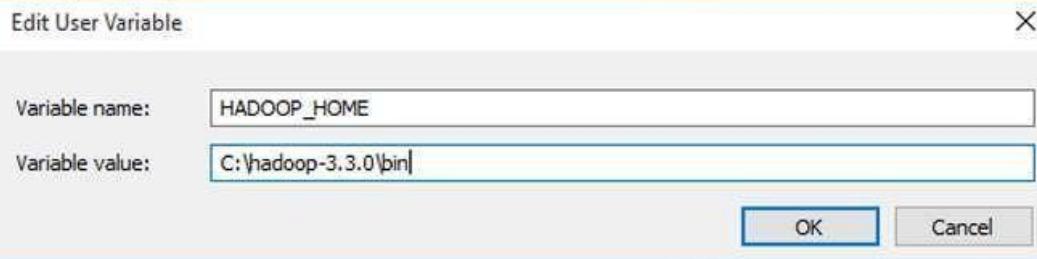
2. Download Hadoop

<https://www.apache.org/dyn/closer.cgi/hadoop/common/hadoop-3.3.0/hadoop-3.3.0.tar.gz>
 right click .rar.gz file -> show more options -> 7-zip->and extract to C:\Hadoop- 3.3.0\

**3. Set the path JAVA_HOME Environment variable****4. Set the path HADOOP_HOME Environment variable**



Click on New to both user variables and system variables.



Click on user variable -> path -> edit-> add path for Hadoop and java upto “bin”

Click Ok, Ok, Ok.

5. Configurations

Edit file C:/Hadoop-3.3.0/etc/hadoop/core-site.xml, paste the xml code in folder and save
<configuration>

```
<property>
<name>fs.defaultFS</name>
<value>hdfs://localhost:9000</value>
</property>
</configuration>
```

Rename “mapred-site.xml.template” to “mapred-site.xml” and edit this file

C:/Hadoop3.3.0/etc/hadoop/mapred-site.xml, paste xml code and save this file.

```
<configuration>
<property>
<name>mapreduce.framework.name</name>
<value>yarn</value>
</property>
</configuration>
```

Create folder “data” under “C:\Hadoop-3.3.0”

Create folder “datanode” under “C:\Hadoop-3.3.0\data”

Create folder “namenode” under “C:\Hadoop-3.3.0\data”

Edit file C:\Hadoop-3.3.0/etc/hadoop/hdfs-site.xml, paste xml code and save this file.

```
<configuration>
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
<property>
<name>dfs.namenode.name.dir</name>
<value>/hadoop-3.3.0/data/namenode</value>
</property>
<property>
<name>dfs.datanode.data.dir</name>
<value>/hadoop-3.3.0/data/datanode</value>
</property>
</configuration>
```

Edit file C:/Hadoop-3.3.0/etc/hadoop/yarn-site.xml, paste xml code and save this file.

```
<configuration>
<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>
<property>
<name>yarn.nodemanager.auxservices.mapreduce.shuffle.class</name>
<value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
<property>
<name>yarn.resourcemanager.address</name>
<value>127.0.0.1:8032</value>
<name>yarn.resourcemanager.scheduler.address</name>
<value>127.0.0.1:8030</value>
<name>yarn.resourcemanager.resource-tracker.address</name>
<value>127.0.0.1:8031</value>
</property>
</configuration>
```

6. Edit file C:/Hadoop-3.3.0/etc/hadoop/hadoop-env.cmd

Find “JAVA_HOME=%JAVA_HOME%” and replace it as

```
set JAVA_HOME="C:\Java\jdk1.8.0_361"
```

7. Download “redistributable” package

Download and run VC_redist.x64.exe

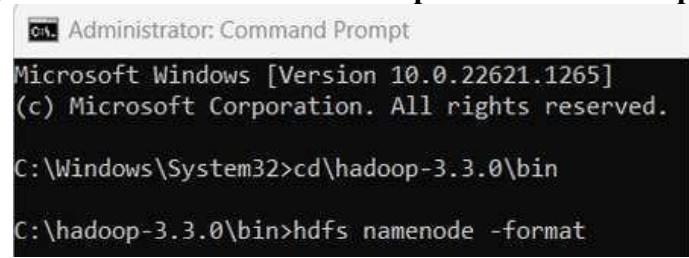
This is a “redistributable” package of the Visual C runtime code for 64-bit applications, from Microsoft. It contains certain shared code that every application written with Visual C expects to have available on the Windows computer it runs on.

8. Hadoop Configurations

Download bin folder from <https://github.com/s911415/apache-hadoop-3.1.0-winutils>

- Copy the bin folder to c:\hadoop-3.3.0. Replace the existing bin folder.

9. copy "hadoop-yarn-server-timelineservice-3.0.3.jar" from ~\hadoop-3.0.3\share\hadoop\yarn\timelineservice to ~\hadoop-3.0.3\share\hadoop\yarn folder.



```
Administrator: Command Prompt
Microsoft Windows [Version 10.0.22621.1265]
(c) Microsoft Corporation. All rights reserved.

C:\Windows\System32>cd\hadoop-3.3.0\bin

C:\hadoop-3.3.0\bin>hdfs namenode -format

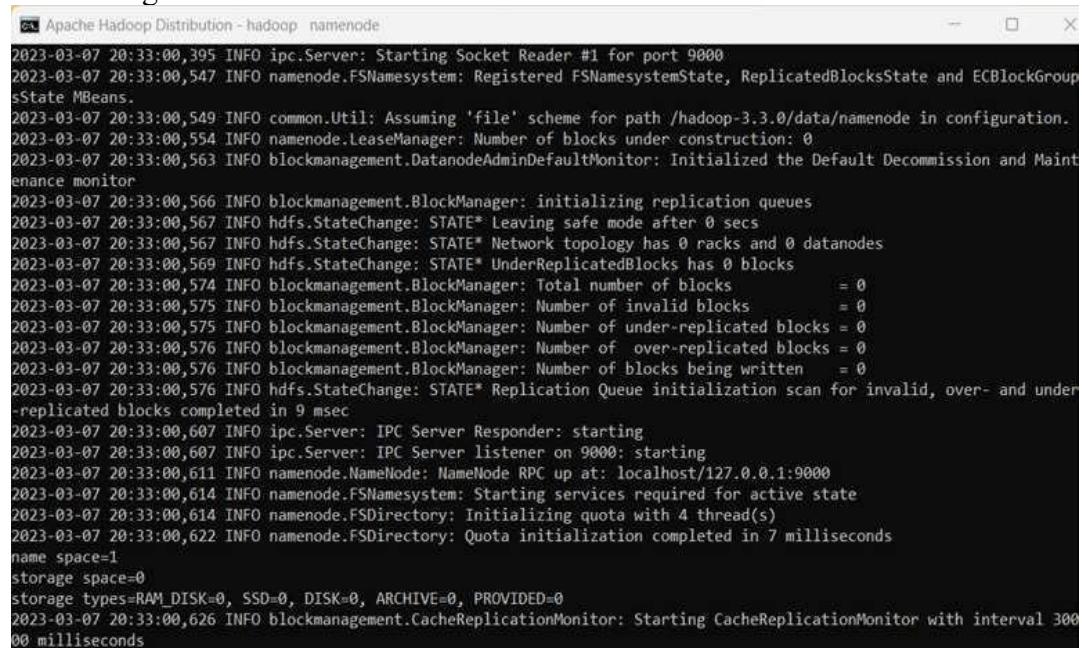
2023-03-07 21:31:34,685 INFO namenode.FSImageFormatProtobuf: Saving image file C:\hadoop-3.3.0\data\namenode\current\fsimage.ckpt_00000000000000000000 using no compression
2023-03-07 21:31:34,844 INFO namenode.FSImageFormatProtobuf: Image file C:\hadoop-3.3.0\data\namenode\current\fsimage.ckpt_00000000000000000000 of size 400 bytes saved in 0 seconds .
2023-03-07 21:31:34,860 INFO namenode.NNStorageRetentionManager: Going to retain 1 images with txid >= 0
2023-03-07 21:31:34,869 INFO namenode.FSImage: FSImageSaver clean checkpoint: txid=0 when meet shutdown.
2023-03-07 21:31:34,870 INFO namenode.NameNode: SHUTDOWN_MSG:
/************************************************************
SHUTDOWN_MSG: Shutting down NameNode at DESKTOP-OL8EULH/192.168.1.19
************************************************************/
```

10. Format the NameNode

- Open cmd „Run as Administrator“ and type command “hdfs namenode –format”

11. Testing

- Open cmd „Run as Administrator“ and change directory to C:\Hadoop-3.3.0\sbin
- type start-all.cmd OR type start-dfs.cmd
- type start-yarn.cmd
- You will get 4 more running threads for Datanode, namenode, resource manager and node manager



```
Apache Hadoop Distribution - hadoop namenode
2023-03-07 20:33:00,395 INFO ipc.Server: Starting Socket Reader #1 for port 9000
2023-03-07 20:33:00,547 INFO namenode.FSNamesystem: Registered FSNamesystemState, ReplicatedBlocksState and ECBLOCKGroupState MBeans.
2023-03-07 20:33:00,549 INFO common.Util: Assuming 'file' scheme for path /hadoop-3.3.0/data/namenode in configuration.
2023-03-07 20:33:00,554 INFO namenode.LeaseManager: Number of blocks under construction: 0
2023-03-07 20:33:00,563 INFO blockmanagement.DatanodeAdminDefaultMonitor: Initialized the Default Decommission and Maintenance monitor
2023-03-07 20:33:00,566 INFO blockmanagement.BlockManager: initializing replication queues
2023-03-07 20:33:00,567 INFO hdfs.StateChange: STATE* Leaving safe mode after 0 secs
2023-03-07 20:33:00,567 INFO hdfs.StateChange: STATE* Network topology has 0 racks and 0 datanodes
2023-03-07 20:33:00,569 INFO hdfs.StateChange: STATE* UnderReplicatedBlocks has 0 blocks
2023-03-07 20:33:00,574 INFO blockmanagement.BlockManager: Total number of blocks = 0
2023-03-07 20:33:00,575 INFO blockmanagement.BlockManager: Number of invalid blocks = 0
2023-03-07 20:33:00,575 INFO blockmanagement.BlockManager: Number of under-replicated blocks = 0
2023-03-07 20:33:00,576 INFO blockmanagement.BlockManager: Number of over-replicated blocks = 0
2023-03-07 20:33:00,576 INFO blockmanagement.BlockManager: Number of blocks being written = 0
2023-03-07 20:33:00,576 INFO hdfs.StateChange: STATE* Replication Queue initialization scan for invalid, over- and under-replicated blocks completed in 9 msec
2023-03-07 20:33:00,607 INFO ipc.Server: IPC Server Responder: starting
2023-03-07 20:33:00,607 INFO ipc.Server: IPC Server listener on 9000: starting
2023-03-07 20:33:00,611 INFO namenode.NameNode: NameNode RPC up at: localhost/127.0.0.1:9000
2023-03-07 20:33:00,614 INFO namenode.FSNamesystem: Starting services required for active state
2023-03-07 20:33:00,614 INFO namenode.FSDirectory: Initializing quota with 4 thread(s)
2023-03-07 20:33:00,622 INFO namenode.FSDirectory: Quota initialization completed in 7 milliseconds
name space=1
storage space=0
storage types=RAM_DISK=0, SSD=0, DISK=0, ARCHIVE=0, PROVIDED=0
2023-03-07 20:33:00,626 INFO blockmanagement.CacheReplicationMonitor: Starting CacheReplicationMonitor with interval 30000 milliseconds
```

Output :

```
C:\hadoop-3.3.0\sbin>jps
5632 Jps
7572 DataNode
3752 ResourceManager
7992 NameNode
8028 NodeManager
```

10. Type JPS command to start-all.cmd command prompt, you will get following output.
11. Run <http://localhost:9870/> from any browser

The screenshot shows the HDFS Health Overview page at <http://localhost:9870/dfshealth.html#tab-overview>. The top navigation bar includes links for Hadoop, Overview, Datanodes, Datanode Volume Failures, Snapshot, Startup Progress, and Utilities. The main content area displays the Overview for 'localhost:9000' (✓active), listing cluster statistics such as Started (Wed Mar 15 12:10:54 +0530 2023), Version (3.3.0, raa96f1871bfd858f9bac59cf2a81ec470da649af), Compiled (Tue Jul 07 00:14:00 +0530 2020 by brahma from branch-3.3.0), Cluster ID (CID-1986aba8-0ed3-43a2-9db7-42944ec51802), and Block Pool ID (BP-1049749432-192.168.56.1-1678862097216). Below this is a 'Browse Directory' section with a table header for columns like Permission, Owner, Group, Size, Last Modified, Replication, Block Size, and Name, though no data is present.

PRACTICAL NO: 06**AIM: Implement word count / frequency programs using MapReduce****Steps:**

- C:\hadoop-3.3.0\sbin>start-dfs.cmd
- C:\hadoop-3.3.0\sbin>start-yarn.cmd
- Open a command prompt as administrator and run the following command to create an input and output folder on the Hadoop file system, to which we will be moving the sample.txt file for our analysis.
- C:\hadoop-3.3.0\bin>cd\
- C:\>hadoop dfsadmin -safemode leave
- DEPRECATED: Use of this script to execute hdfs command is deprecated.
- Instead use the hdfs command for it.
- Safe mode is OFF
- C:\>hadoop fs -mkdir /input_dir
- Check it by giving the following URL at browser http://localhost:9870
- Now apply the following command at c:\>
- C:\>hadoop fs -put C:/input_file.txt /input_dir
- Verify input_file.txt available in HDFS input directory (input_dir).
- C:\>Hadoop fs -ls /input_dir
- C:\>hadoop jar C:/hadoop-3.3.0/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.3.0.jar wordcount /input_dir /output_dir
- In case, there is some error in executing then copy the file MapReduceClient.jar in C:\ and run the program with the jar file using existing MapReduceClient.jar file as:
- C:\> hadoop jar C:/MapReduceClient.jar wordcount /input_dir /output_dir
- Now, check the output_dir on browser as follows:
- Click on output_dir à part-r-00000 à Head the file (first 32 K) and check the file content as the output.
- Alternatively, you may type the following command on CMD window as:
- C:\>hadoop dfs -cat /output_dir/*

You can get the following output

Output:

```

hadoop dfs -cat /output_dir/*
Hadoop 1
Window 1
version 1
is 1
easy 1
compared 1
to 1
Ubuntu 1

```

PRACTICAL NO: 07**AIM: Implement Decision tree classification techniques**

A **Decision Tree** is one of the most intuitive classification techniques because it mimics human decision-making. We will use the `rpart` package to build the tree and `rpart.plot` to visualize it.

For this example, we will use the built-in `iris` dataset, which classifies flowers based on their physical measurements.

Step 1: Install and Load Required Packages

We need `rpart` for the logic and `rpart.plot` for a clean, easy-to-read visualization.

How to execute: Copy these lines into your R script and run them.

R

```
# Install relevant packages
install.packages("rpart")
install.packages("rpart.plot")
```

Load the libraries

```
library(rpart)
library(rpart.plot)
```

Step 2: Prepare and Split the Data

We split the data to ensure our model can generalize to new, unseen information.

R

```
# Load data
data(iris)
```

Set seed for reproducibility

```
set.seed(42)
```

Split data: 70% Training, 30% Testing

```
sample_index <- sample(1:nrow(iris), 0.7 * nrow(iris))
train_data <- iris[sample_index, ]
test_data <- iris[-sample_index, ]
```

Step 3: Build the Decision Tree Classifier

We use the `rpart()` function. The formula `Species ~ .` tells R to predict the "Species" using all other available variables (Sepal Length, Width, etc.).

R

```
# Train the Decision Tree model
dt_model <- rpart(Species ~ ., data = train_data, method = "class")

# View a text summary of the tree
print(dt_model)
```

Step 4: Visualize the Tree

The best part of a Decision Tree is the visualization. This makes it "Explainable AI."

R

```
# Plot the tree
rpart.plot(dt_model, type = 4, extra = 101, under = TRUE, cex = 0.8, box.palette = "GnBu")
```

Step 5: Evaluate the Performance

Finally, we test the model on the 30% of data we held back to see its accuracy.

R

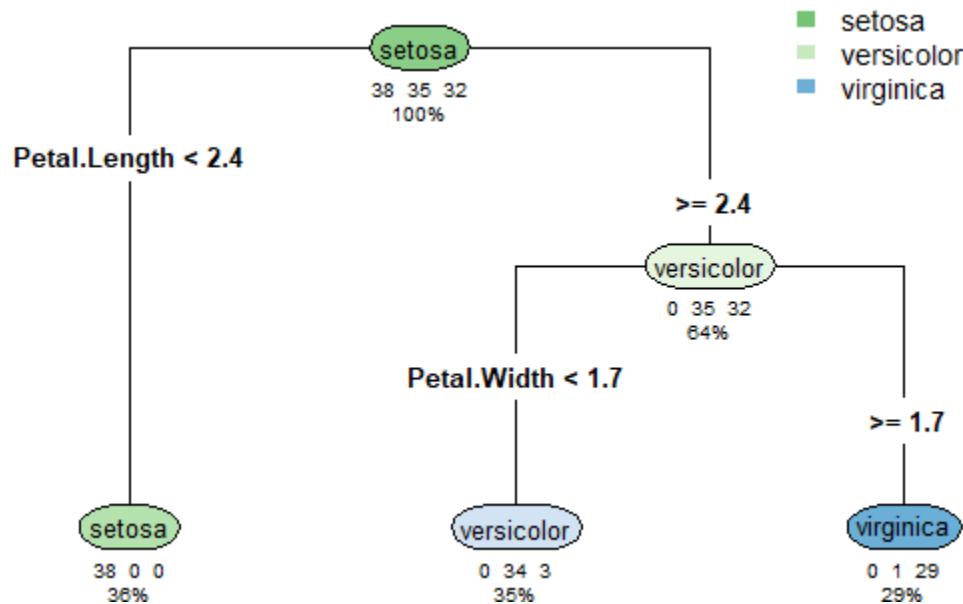
```
# Make predictions on the test set
dt_predictions <- predict(dt_model, test_data, type = "class")
```

```
# Create a Confusion Matrix
conf_matrix <- table(test_data$Species, dt_predictions)
print(conf_matrix)
```

```
# Calculate Accuracy
accuracy <- sum(diag(conf_matrix)) / sum(conf_matrix)
print(paste("Accuracy:", round(accuracy, 4)))
```

OUTPUT:-

```
print(conf_matrix)
      dt_predictions
      setosa versicolor virginica
setosa       12          0          0
versicolor     0         14          1
virginica      0          1         17
```



```
print(paste("Accuracy:", round(accuracy, 4)))
[1] "Accuracy: 0.9556"
```

PRACTICAL NO: 08**AIM: Implement SVM classification techniques**

Support Vector Machines (SVM) are powerful supervised learning models used for classification. SVM works by finding the **hyperplane** (a boundary) that best separates different classes with the maximum possible margin.

To implement SVM in R, the most popular and reliable package is `e1071`.

1. Required Packages

In R, SVM is commonly implemented using the `e1071` package.

```
install.packages("e1071")      # run once
library(e1071)
```

2. Load Dataset (Iris Dataset)

```
data(iris)
head(iris)
```

- **Features:** Sepal.Length, Sepal.Width, Petal.Length, Petal.Width
- **Target/Class:** Species

3. Split Dataset into Training & Testing

```
set.seed(123)

index <- sample(1:nrow(iris), 0.7 * nrow(iris))
train_data <- iris[index, ]
test_data <- iris[-index, ]
```

4. Linear SVM Classification

```
svm_linear <- svm(
  Species ~ .,
  data = train_data,
  kernel = "linear"
)

pred_linear <- predict(svm_linear, test_data)

table(Predicted = pred_linear, Actual = test_data$Species)
mean(pred_linear == test_data$Species)
```

Confusion Matrix

```
library(caret)

confusionMatrix(pred_rbf, test_data$Species)
```

OUTPUT:

```
pred_linear <- predict(svm_linear, test_data)
> table(Predicted = pred_linear, Actual = test_data$Species)
    Actual
Predicted   setosa versicolor virginica
  setosa      14        0        0
  versicolor    0       17        0
  virginica     0        1       13
> mean(pred_linear == test_data$Species)
[1] 0.9777778
```

- **Rows = Predicted class**
- **Columns = Actual class**

Correct Predictions (Diagonal Values)

These are the **correct classifications**:

- **Setosa:** 14 correctly classified
- **Versicolor:** 17 correctly classified
- **Virginica:** 13 correctly classified

Incorrect Predictions (Off-diagonal Values)

- **1 Versicolor was misclassified as Virginica**

All other samples are classified correct